

## Naive Bayes Classifier

(Code: Subhajit Das)

### What is Naive Bayes Classifier:

The Naive Bayes Classifier is a supervised machine learning algorithm used for classification tasks<sup>1</sup>. It's part of a family of generative learning algorithms, which means it seeks to model the distribution of inputs of a given class or category.

Naive Bayes classifiers are a collection of classification algorithms based on Bayes' Theorem. They are not a single algorithm but a family of algorithms where all of them share a common principle: every pair of features being classified is independent of each other. This means that the algorithm assumes that no pair of features are dependent, and each feature contributes equally to the outcome.

$$P(A | B) = \frac{P(B | A)P(A)}{P(B)}$$

where:

$P(A|B)$  = Conditional Probability of A given B

$P(B|A)$  = Conditional Probability of A given B

$P(A)$  = Probability of event A

$P(B)$  = Probability of event A

### Where we can use Naive Bayes Classifier:

The Naive Bayes classifier can be used in a variety of applications, including:

1. **Spam Filtering:** It's commonly used in email services to determine whether an email is spam or not.
2. **Text Classification:** It can be used to categorize documents into different categories.
3. **Sentiment Analysis:** It's used to analyze the sentiment of a given text, such as determining whether a movie review is positive or negative.
4. **Recommender Systems:** It can be used to recommend products or services to users based on their past behavior.
5. **Credit Scoring:** It's used in the finance industry to determine the creditworthiness of a customer.
6. **Medical Data Classification:** It's used in the healthcare industry to classify patient data into different categories for better diagnosis and treatment.

### How Naive Bayes Classifier works:

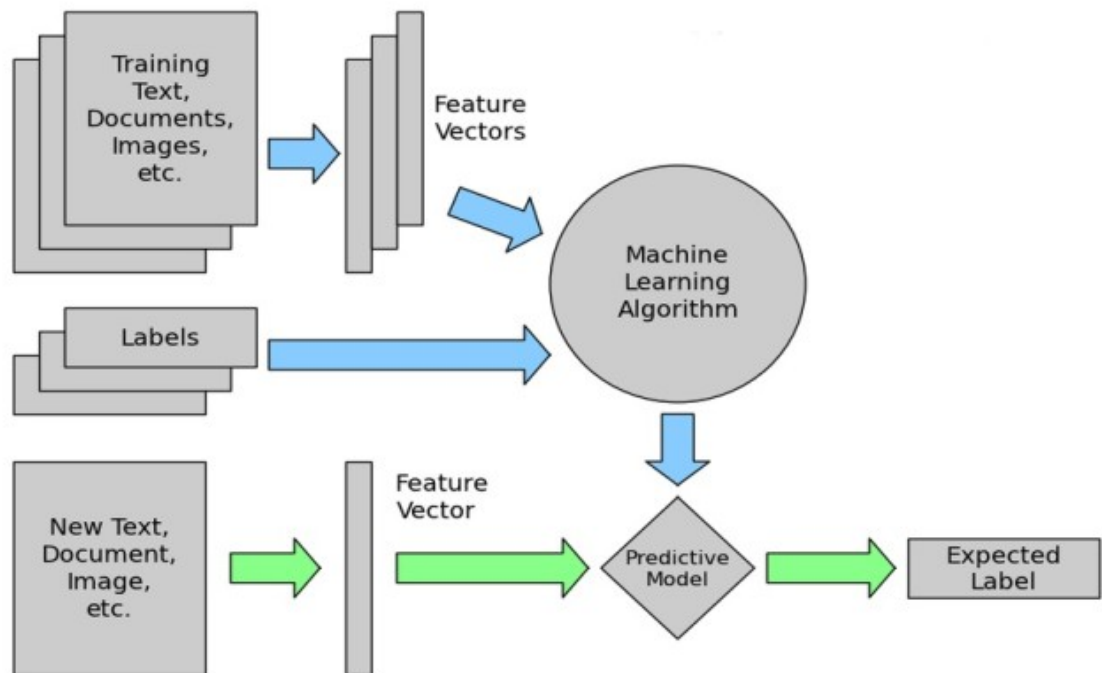
The Naive Bayes classifier works based on the principles of probability theory and Bayes' Theorem. Here's a step-by-step explanation of how it works:

1. **Feature Independence:** The fundamental assumption of Naive Bayes is that each feature makes an independent and equal contribution to the outcome. This means that the algorithm assumes that no pair of features are dependent.
2. **Dataset Division:** The dataset is divided into a feature matrix and a response vector. The feature matrix contains all the vectors (rows) of the dataset, where each vector consists of the value of dependent features. The response vector contains the value of

the class variable (prediction or output) for each row of the feature matrix.

3. **Application of Bayes' Theorem:** Bayes' Theorem is used to find the probability of an event occurring given the probability of another event that has already occurred. In the context of Naive Bayes, we are trying to find the probability of a class (or tag), given the features of the data.
4. **Prediction:** The Naive Bayes classifier operates by returning the class which has the maximum posterior probability out of a group of classes.

Despite its simplicity and the 'naive' assumption of feature independence, the Naive Bayes classifier often performs well and is widely used in various applications.



```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
```

```
In [2]: ads_df = pd.read_csv("/content/drive/MyDrive/ML and DL DataSets/8_Social_Net
ads_df.head()
```

```
Out[2]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

```
In [3]: ads_df.info
```

```
Out[3]: <bound method DataFrame.info of
y  Purchased
0    15624510    Male    19          19000          0
1    15810944    Male    35          20000          0
2    15668575  Female    26          43000          0
3    15603246  Female    27          57000          0
4    15804002    Male    19          76000          0
..      ...      ...      ...          ...          ...
395  15691863  Female    46          41000          1
396  15706071    Male    51          23000          1
397  15654296  Female    50          20000          1
398  15755018    Male    36          33000          0
399  15594041  Female    49          36000          1

[400 rows x 5 columns]>
```

### Separating features and labels

```
In [4]: x = ads_df.iloc[:, [1,2,3]].values # '.values' for Numpy representation of t
x
```

```
Out[4]: array(['Male', 19, 19000],
              ['Male', 35, 20000],
              ['Female', 26, 43000],
              ...,
              ['Female', 50, 20000],
              ['Male', 36, 33000],
              ['Female', 49, 36000]), dtype=object)
```

```
In [5]: y = ads_df.iloc[:, 4].values
y
```

```
Out[5]: array([0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1,
              1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
              0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
              0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
              0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
              0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
              0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
              0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
              0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1,
              0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0,
              1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1,
              0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 0, 0, 0, 1,
              1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1,
              0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0,
              1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1,
              0, 1, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1,
              1, 1, 0, 1])
```

### Using LabelEncoder for Gender column

```
In [6]: from sklearn.preprocessing import LabelEncoder
```

```
In [7]: Le = LabelEncoder()
```

```
In [8]: x[:,0] = Le.fit_transform(x[:,0])
```

```
In [9]: x # Gender converted through Label Encoder
```

```
Out[9]: array([[1, 19, 19000],
               [1, 35, 20000],
               [0, 26, 43000],
               ...,
               [0, 50, 20000],
               [1, 36, 33000],
               [0, 49, 36000]], dtype=object)
```

## Analysis

```
In [10]: print(ads_df['Gender'].value_counts()) # Total gender of Male and Female
```

```
Female    204
Male       196
Name: Gender, dtype: int64
```

```
In [11]: print(ads_df['Purchased'].value_counts()) # Total purchased and not purchased
```

```
0    257
1    143
Name: Purchased, dtype: int64
```

```
In [12]: print(ads_df.groupby('Gender')['Purchased'].value_counts()) # Purchasing by
```

```
Gender  Purchased
Female  0           127
        1           77
Male    0           130
        1           66
Name: Purchased, dtype: int64
```

## Splitting train and test datasets

```
In [13]: from sklearn.model_selection import train_test_split
```

```
In [14]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2)
```

```
In [15]: # Length or sample of train dataset
len(x_train)
```

```
Out[15]: 320
```

```
In [16]: # Length or sample of test dataset
len(x_test)
```

```
Out[16]: 80
```

## Using Naive Bayes Classifier

### Parameters used in GaussianNB():

The GaussianNB in sklearn has the following parameters:

1. **priors**: Prior probabilities of the classes. If specified, the priors are not adjusted according to the data.
2. **var\_smoothing**: Portion of the largest variance of all features that is added to variances for calculation stability.

These parameters allow you to control the behavior of the Gaussian Naive Bayes classifier and can be tuned to improve the performance of the model on your specific task. For example, you can tune the 'var\_smoothing' parameter

```
In [17]: from sklearn.naive_bayes import GaussianNB
```

```
In [18]: naive = GaussianNB()
```

```
In [19]: naive.fit(x_train, y_train)
```

```
Out[19]: GaussianNB()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [20]: naive.predict(x_test) # 1: purchased, 0: not purchased
```

```
Out[20]: array([1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
                0, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0,
                1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0])
```

```
In [21]: x_test[:5]# 1: male, 0: Female
```

```
Out[21]: array([[0, 48, 119000],
                [0, 32, 117000],
                [1, 40, 61000],
                [0, 48, 96000],
                [0, 30, 116000]], dtype=object)
```

```
In [22]: y_test[:5]
```

```
Out[22]: array([1, 1, 0, 1, 0])
```

### Viewing the prediction score

```
In [23]: naive.score(x_test, y_test)
```

```
Out[23]: 0.8875
```

```
In [24]: from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
```

```
In [34]: pred_naive = naive.predict(x_test) # storing the predictions that your model
```

```
In [26]: # Classification metrics can't handle a mix of multiclass-multioutput and multiclass-multiclass
confusion_matrix(y_test, pred_naive)
```

```
Out[26]: array([[47,  4],
                [ 5, 24]])
```

```
In [27]: print(accuracy_score(y_test, pred_naive))
```

0.8875

```
In [28]: print(classification_report(y_test, pred_naive))
```

	precision	recall	f1-score	support
0	0.90	0.92	0.91	51
1	0.86	0.83	0.84	29
accuracy			0.89	80
macro avg	0.88	0.87	0.88	80
weighted avg	0.89	0.89	0.89	80

### Predict purchased with user input

```
In [29]: enc_gender = LabelEncoder()  
categories = ['Male', 'Female']  
enc_gender.fit(categories)
```

Out[29]: LabelEncoder()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**  
**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [30]: purchased_le = LabelEncoder()  
categories = ['yes', 'no']  
purchased_le.fit(categories)
```

Out[30]: LabelEncoder()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**  
**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [31]: gender = input("Enter the Gender: ") # 0: 'Female', 1: 'male'  
gender_var = enc_gender.transform([gender])[0] # The [0] at the end is an index  
  
age = int(input("Enter the Age: "))  
  
estimatedsalary = int(input("Enter the Estimated Salary: "))  
  
label_map = {0: 'no', 1: 'yes'}  
predict_value = naive.predict([[gender_var, age, estimatedsalary]])  
  
# Map the numerical prediction back to a string label  
label_value = label_map[predict_value[0]]  
  
print(label_value) # 0: 'no', 1: 'yes'
```

Enter the Gender: Male  
Enter the Age: 21  
Enter the Estimated Salary: 18000  
no

## Plotting on actual vs predicted data

```
In [32]: import matplotlib.pyplot as plt

# Assuming y_actual and y_predicted are your data
y_actual = naive.predict(x_test)
y_predicted = y_test

plt.figure(figsize=(10, 6))

# Plotting the actual values
plt.scatter(range(len(y_actual)), y_actual, color='blue', label='Actual')

# Plotting the predicted values
plt.scatter(range(len(y_predicted)), y_predicted, color='red', label='Predicted')

plt.title('Actual vs Predicted')
plt.xlabel('Index')
plt.ylabel('Value')
plt.legend(loc='upper left')
plt.show()
```

